

红酒评论数据集

1. 数据集加载

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import pandas as pd
4 %matplotlib inline
5
6 path=['./dataset/wine_reviews/winemag-data-130k-v2.csv',
7       './dataset/wine_reviews/winemag-data_first150k.csv']
8 data = pd.concat([pd.read_csv(p) for p in path])
9 data.info()
```

```
1 <class 'pandas.core.frame.DataFrame'>
2 Int64Index: 280901 entries, 0 to 150929
3 Data columns (total 14 columns):
4 #   Column           Non-Null Count  Dtype  
5 ---  --  
6 0   Unnamed: 0        280901 non-null  int64  
7 1   country          280833 non-null  object  
8 2   description      280901 non-null  object  
9 3   designation      197701 non-null  object  
10 4   points           280901 non-null  int64  
11 5   price            258210 non-null  float64 
12 6   province         280833 non-null  object  
13 7   region_1         234594 non-null  object  
14 8   region_2         111464 non-null  object  
15 9   taster_name      103727 non-null  object  
16 10  taster_twitter_handle 98758 non-null  object  
17 11  title            129971 non-null  object  
18 12  variety          280900 non-null  object  
19 13  winery           280901 non-null  object  
20 dtypes: float64(1), int64(2), object(11)
21 memory usage: 32.1+ MB
```

2. 数据集可视化可摘要

2.1 数据摘要和可视化

- 共13个属性(ID除外),其中:
- 标称属性:

1. country 国家
2. designation 葡萄园
3. province 省份
4. region_1 葡萄酒生产地
5. region_2 详细生产地
6. variety 葡萄酒种类
7. winery 酿酒厂

- 数值属性:

8. points 评分

9. price 价格

- 其他:

10. title 评价标题

11. description 评价内容

12. taster_name 评价人


13. taster_twitter_handle 评价人twitter

(1) country 属性

```
1 | print(data['country'].value_counts(dropna = False).head(10))
2 | data['country'].value_counts(dropna = False).plot(kind="bar", figsize=(20, 4))
```

```
1 | US           116901
2 | France       43191
3 | Italy         43018
4 | Spain         14913
5 | Portugal      11013
6 | Chile          10288
7 | Argentina     9431
8 | Australia      7286
9 | Austria        6402
10 | New Zealand    4739
11 | Name: country, dtype: int64
```

```
1 | <AxesSubplot:>
```



(2) designation 属性


```
1 | print(data['designation'].value_counts(dropna = False).head(10))
2 | data['designation'].value_counts(dropna = False)
   [:100].plot(kind="bar", figsize=(20, 4))
```

```

1  Nan          83200
2  Reserve      4761
3  Reserva      3069
4  Estate        2893
5  Barrel sample 1701
6  Riserva       1452
7  Brut          1137
8  Estate Grown 1070
9  Barrel Sample 891
10 Crianza       846
11 Name: designation, dtype: int64

```

```
1 | <AxesSubplot:>
```



(3) province 属性

```

1 | print(data['province'].value_counts(dropna = False).head(10))
2 | data['province'].value_counts(dropna = False)[:100].plot(kind="bar", figsize=(20,4))


```

```

1 | California      80755
2 | Washington      18389
3 | Tuscany         13178
4 | Bordeaux        12052
5 | Oregon          9962
6 | Northern Spain   8743
7 | Burgundy        8288
8 | Mendoza Province 8006
9 | Piedmont        7822
10 | Veneto          6678
11 | Name: province, dtype: int64

```

```
1 | <AxesSubplot:>
```




(4) region_1 属性

```
1 | print(data['region_1'].value_counts(dropna = False).head(10))
2 | data['region_1'].value_counts(dropna = False)[:100].plot(kind="bar", figsize=(20,4))
```

1	NaN	46307
2	Napa Valley	10689
3	Columbia Valley (WA)	9099
4	Russian River Valley	6662
5	California	6091
6	Mendoza	5887
7	Paso Robles	5403
8	Willamette Valley	4397
9	Alsace	3737
10	Rioja	3362
11	Name: region_1, dtype:	int64

1 | <AxesSubplot:>




(5) region_2 属性

```
1 | print(data['region_2'].value_counts(dropna = False).head(10))
2 | data['region_2'].value_counts(dropna = False).plot(kind="bar", figsize=(20,4))
```

```
1 |   Nan          169437
2 | Central Coast    24122
3 | Sonoma         20286
4 | Columbia Valley 17260
5 | Napa           15615
6 | Willamette Valley 6604
7 | California Other 6179
8 | Finger Lakes    3287
9 | Sierra Foothills 3122
10 | Napa-Sonoma    2814
11 | Name: region_2, dtype: int64
```

```
1 | <AxesSubplot:>
```




(6) variety 属性

```
1 | print(data['variety'].value_counts(dropna = False).head(10))
2 | data['variety'].value_counts(dropna = False)[:100].plot(kind="bar", figsize=(20,4))
```

```
1 | Pinot Noir      27563
2 | Chardonnay      26235
3 | Cabernet Sauvignon 22272
4 | Red Blend       19008
5 | Bordeaux-style Red Blend 14262
6 | Sauvignon Blanc 11287
7 | Riesling         10713
8 | Syrah            9967
9 | Merlot           8172
10 | zinfandel       6513
11 | Name: variety, dtype: int64
```

```
1 | <AxesSubplot:>
```




(7) winery 属性

```
1 print(data['winery'].value_counts(dropna = False).head(10))
2 data['winery'].value_counts(dropna = False)[:100].plot(kind="bar", figsize=(20,4))
```

1	williams selyem	585
2	Testarossa	492
3	DFJ Vinhos	473
4	Chateau Ste. Michelle	419
5	wines & winemakers	391
6	Georges Duboeuf	384
7	Concha y Toro	380
8	Columbia Crest	376
9	Louis Latour	353
10	Kendall-Jackson	346
11	Name: winery, dtype: int64	

1 | <AxesSubplot:>




(8) points 属性

```
1 print('最小值', data['points'].min())
2 print('第一四分位数', data['points'].quantile(0.25))
3 print('中位数', data['points'].median())
4 print('第三四分位数', data['points'].quantile(0.75))
5 print('最大值', data['points'].max())
6 print('缺失值个数', data['points'].isnull().sum())
```

```
1 | 最小值 80
2 | 第一四分位数 86.0
3 | 中位数 88.0
4 | 第三四分位数 90.0
5 | 最大值 100
6 | 缺失值个数 0
```

```
1 | data.boxplot(['points'],return_type='dict')
```

```
1 | {'whiskers': [<matplotlib.lines.Line2D at 0x2b839845fc8>,
2 |   <matplotlib.lines.Line2D at 0x2b8398614c8>],
3 |   'caps': [<matplotlib.lines.Line2D at 0x2b839868908>,
4 |   <matplotlib.lines.Line2D at 0x2b839868d48>],
5 |   'boxes': [<matplotlib.lines.Line2D at 0x2b839861408>],
6 |   'medians': [<matplotlib.lines.Line2D at 0x2b839868d88>],
7 |   'fliers': [<matplotlib.lines.Line2D at 0x2b839868e48>],
8 |   'means': []}
```



```
1 | def detect_outliers(sr):
2 |     q1 = sr.quantile(0.25)
3 |     q3 = sr.quantile(0.75)
4 |     iqr = q3-q1 #Interquartile range
5 |     fence_low = q1-1.5*iqr
6 |     fence_high = q3+1.5*iqr
7 |     outliers = sr.loc[(sr < fence_low) | (sr > fence_high)]
8 |     return outliers
9 | detect_outliers(data['points'])
10 |
```

```
1 345      100
2 346      98
3 347      97
4 348      97
5 349      97
6 ...
7 143626   97
8 144494   97
9 144652   97
10 148047  97
11 149172  97
12 Name: points, Length: 928, dtype: int64
```

其中points>97的均视为离群点


(9) price 属性

```
1 print('最小值',data['price'].min())
2 print('第一四分位数',data['price'].quantile(0.25))
3 print('中位数',data['price'].median())
4 print('第三四分位数',data['price'].quantile(0.75))
5 print('最大值',data['price'].max())
6 print('缺失值个数',data['price'].isnull().sum())
```

```
1 最小值 4.0
2 第一四分位数 16.0
3 中位数 25.0
4 第三四分位数 40.0
5 最大值 3300.0
6 缺失值个数 22691
```

```
1 data.boxplot(['price'],return_type='dict')
```

```
1 {'whiskers': [<matplotlib.lines.Line2D at 0x1c667b47d08>,
2                 <matplotlib.lines.Line2D at 0x1c667b36ac8>],
3  'caps': [<matplotlib.lines.Line2D at 0x1c667b36a48>,
4            <matplotlib.lines.Line2D at 0x1c667b341c8>],
5  'boxes': [<matplotlib.lines.Line2D at 0x1c667b44388>],
6  'medians': [<matplotlib.lines.Line2D at 0x1c667b34d08>],
7  'fliers': [<matplotlib.lines.Line2D at 0x1c667b32308>],
8  'means': []}
```



```

1 | def detect_outliers(sr):
2 |     q1 = sr.quantile(0.25)
3 |     q3 = sr.quantile(0.75)
4 |     iqr = q3-q1 #Interquartile range
5 |     fence_low  = q1-1.5*iqr
6 |     fence_high = q3+1.5*iqr
7 |     outliers = sr.loc[(sr < fence_low) | (sr > fence_high)]
8 |     return outliers
9 | detect_outliers(data['price'])

```

```

1 |   60      100.0
2 | 111      85.0
3 | 118      80.0
4 | 119      80.0
5 | 134      78.0
6 | ...
7 | 150570    90.0
8 | 150613    80.0
9 | 150727    83.0
10 | 150762    100.0
11 | 150765    87.0
12 | Name: price, Length: 15128, dtype: float64

```

3. 数据缺失的处理

统计所有属性的数据缺失情况：

```
1 | print(data.isnull().sum(axis=0))
```

1	Unnamed: 0	0
2	country	68
3	description	0
4	designation	83200
5	points	0

```

6 price 22691
7 province 68
8 region_1 46307
9 region_2 169437
10 taster_name 177174
11 taster_twitter_handle 182143
12 title 150930
13 variety 1
14 winery 0
15 dtype: int64

```

3.1 处理country属性缺失

缺失原因：统计失误，可以通过属性的相关关系来填补缺失值，查找country属性缺失的葡萄园所在国家并进行填充

```

1 import copy
2 designation2country = {
3     "Askitikos": "Greece",
4     "Shah": "US",
5     "Piedra Feliz": "Chile",
6 }
7 data_country = copy.deepcopy(data)
8 for index, row in data_country.iterrows():
9     if not row['country']:
10         row['country'] = designation2country[row['designation']]


```

```

1 import matplotlib.pyplot as plt
2 plt.subplot(2,1,1)
3 data["country"].value_counts(dropna = False).plot(kind='bar', figsize=(20,8))
4 plt.subplot(2,1,2)
5 data_country["country"].value_counts(dropna = False).plot(kind='bar', figsize=(20,8))

```

1 <AxesSubplot:>




3.2 处理designation属性缺失

缺失原因：葡萄园可能较小或葡萄来源不明，将缺失部分剔除

```
1 data_designation = data.dropna(subset=['designation'])
2 plt.subplot(2,1,1)
3 data["designation"].value_counts(dropna = False)
[ :100].plot(kind='bar', figsize=(20,8))
4 plt.subplot(2,1,2)
5 data_designation["designation"].value_counts(dropna = False)
[ :100].plot(kind='bar', figsize=(20,8))
```

```
1 <AxesSubplot:>
```




3.3 处理price属性缺失

缺失原因：通过数据平均值来填补缺失值

```
1 data_price= copy.deepcopy(data)
2 mean_price = data_price['price'].mean()
3 data_price = data_price['price'].fillna(mean_price)
4
5 plt.subplot(2,1,1)
6 plt.boxplot(data_price,vert=False)
7 plt.subplot(2,1,2)
8 data.boxplot(['price'],return_type='dict',vert=False)
```

```
1 {'whiskers': [<matplotlib.lines.Line2D at 0x7feddd199940>,
2 <matplotlib.lines.Line2D at 0x7feddd199c10>],
3 'caps': [<matplotlib.lines.Line2D at 0x7feddd199ee0>,
4 <matplotlib.lines.Line2D at 0x7feddd1371f0>],
5 'boxes': [<matplotlib.lines.Line2D at 0x7feddd199670>],
6 'medians': [<matplotlib.lines.Line2D at 0x7feddd1374c0>],
7 'fliers': [<matplotlib.lines.Line2D at 0x7feddd137790>],
8 'means': []}
```



3.4 处理region_1属性缺失


缺失原因：通过相似对象的最高频率来填补缺失值

```

1 mode = data['region_1'].mode()
2 data_region1 = data['region_1'].fillna(mode)
3
4
5 plt.subplot(2,1,1)
6 data['region_1'].value_counts(dropna = False)[:50].plot(kind='bar', figsize=(20,8))
7 plt.subplot(2,1,2)
8 data_region1.value_counts(dropna = False)[:50].plot(kind='bar', figsize=(20,8))

```

1 <AxesSubplot:>



3.5 处理region_2属性缺失

缺失原因：将缺失部分剔除

```

1 data_region2 = data.dropna(subset=['region_2'])
2 plt.subplot(2,1,1)
3 data["region_2"].value_counts(dropna = False)[:50].plot(kind='bar', figsize=(20,8))
4 plt.subplot(2,1,2)
5 data_region2["region_2"].value_counts(dropna = False)
[:50].plot(kind='bar', figsize=(20,8))

```

1 | <AxesSubplot:>

